

Residual Analysis

Table of Contents

- 1 Model Adequacy Checking
- 2 Raw Residuals
- 3 Standardized Residuals
- 4 Studentized Residuals
- 5 PRESS Residuals
- 6 R-Student Residuals

Table of contents continued in the next page...

Table of Contents (cont.)

- 7 Example: Delivery Time Data
- 8 Normality
- 9 Constant Variance
- 10 Linearity of the Mean Function
- 11 Uncorrelated Errors
- 12 Summary

Model Adequacy Checking

Assumptions Behind the Fitted Model

The linear model $\mathbf{y} = \mathbf{X}\beta + \varepsilon$ is fitted under five assumptions:

- The relationship between y and the regressors is (approximately) linear
- $E[\varepsilon_i] = 0$
- $\text{Var}(\varepsilon_i) = \sigma^2$ for all i (constant variance)
- The errors are uncorrelated
- The errors are normally distributed (needed for the t and F procedures)

Least squares returns estimates whether or not these assumptions hold.

What Can We Observe?

The errors ε_i are never observed. What can we observe instead?

Raw Residuals

Raw Residuals

The i th residual is $e_i = y_i - \hat{y}_i$. In matrix form

$$\begin{aligned}\mathbf{e} &= \mathbf{y} - \hat{\mathbf{y}} = (\mathbf{I} - \mathbf{H})\mathbf{y} \\ &= (\mathbf{I} - \mathbf{H})\boldsymbol{\varepsilon},\end{aligned}$$

where the last step uses $(\mathbf{I} - \mathbf{H})\mathbf{X} = \mathbf{0}$.

Leverage

The hat matrix is $\mathbf{H} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ with diagonal elements

$$h_{ii} = \mathbf{x}_i^\top (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{x}_i.$$

- Intuitively, h_{ii} measures how far $\mathbf{x}_i$ is from the center of the regressor data: points with unusual x -values have more "leverage" to pull the fitted surface toward their own y_i
- h_{ii} is large when $\mathbf{x}_i$ is far from the centroid of the regressor values
- $0 \leq h_{ii} \leq 1$ and $\sum_{i=1}^n h_{ii} = p$, where p is the number of model parameters
- $\hat{y}_i = \sum_j h_{ij} y_j$, so h_{ii} is the weight of y_i in its own fitted value

Flagging High - Leverage Points


Rule of thumb (MPV): observation i has high leverage if

$$h_{ii} > \frac{2p}{n}.$$

- Delivery time data: $p = 3$, $n = 25$, so the cutoff is $2p/n = 2(3)/25 = 0.24$
- Observation 9: $h_{99} = 0.498 > 0.24$, flagged as high leverage
- High leverage is about an unusual $\mathbf{x}_i$, not an unusual y_i : a flagged point can still have a small residual if y_i happens to agree with the trend at that $\mathbf{x}_i$

Leverage: Geometric Intuition

- Leverage depends only on the location of $\mathbf{x}_i$, not on y_i
- A point far from $\bar{x}$ has high leverage, whatever its y -value
- Here the flagged point sits far from the rest of the data in x , yet close to the fitted line: high leverage, small residual



Mean and Covariance of the Residuals

$$\begin{aligned} \mathbf{E}[\mathbf{e}] &= \mathbf{0} \\ \text{Var}(\mathbf{e}) &= \sigma^2(\mathbf{I} - \mathbf{H}) \end{aligned} \tag{1}$$

Element by element

$$\begin{aligned} \text{Var}(e_i) &= \sigma^2(1 - h_{ii}) \\ \text{Cov}(e_i, e_j) &= -\sigma^2 h_{ij}, \quad i \neq j. \end{aligned}$$

Are the Residuals iid?

The errors are assumed iid $N(0, \sigma^2)$. Are the residuals $e_1, \dots, e_n$ iid?

Standardized Residuals

Standardized Residuals

$$d_i = \frac{e_i}{\sqrt{\text{MSE}}}, \quad \text{MSE} = \frac{\text{SS}_{\text{Res}}}{n - p} \quad (2)$$

- $E[d_i] = 0$ and $\text{Var}(d_i) \approx 1$, and d_i is free of the units of y
- Rule of thumb: $|d_i| > 3$ marks a potential outlier
- d_i ignores that $\text{Var}(e_i) = \sigma^2(1 - h_{ii})$ differs across observations, see (1)

Studentized Residuals

Studentized Residuals

$$r_i = \frac{e_i}{\sqrt{\text{MSE}(1 - h_{ii})}} = \frac{d_i}{\sqrt{1 - h_{ii}}} \quad (3)$$

- Divides e_i by its own estimated standard deviation from (1), so $\text{Var}(r_i) = 1$ for every i , whatever the location of $\mathbf{x}_i$
- $|r_i| \geq |d_i|$, with d_i from (2)
- $r_i \approx d_i$ when h_{ii} is small, and they differ noticeably for high-leverage points

Leverage and Residual Size

Points with large leverage tend to have small raw residuals. Why?

PRESS Residuals

Leave-One-Out Prediction Error

The PRESS residual is the error when y_i is predicted from the model fitted to the other $n - 1$ observations

$$e_{(i)} = y_i - \hat{y}_{(i)} = \frac{e_i}{1 - h_{ii}}.$$

- The identity avoids refitting the model n times
- A large h_{ii} inflates the gap between $e_{(i)}$ and e_i

Variance of PRESS Residuals and the PRESS Statistic

$$\text{Var}(e_{(i)}) = \frac{\sigma^2}{1 - h_{ii}}$$

$$\text{PRESS} = \sum_{i=1}^n e_{(i)}^2$$

- PRESS measures predictive ability: smaller is better
- Standardizing $e_{(i)}$ by its standard deviation returns the studentized residual

$$\frac{e_{(i)}}{\sqrt{\text{Var}(e_{(i)})}} = \frac{e_i}{\sqrt{\sigma^2(1 - h_{ii})}}$$

R - Student Residuals

Removing Observation i From the Variance Estimate

MSE uses all n observations, including observation i : a large e_i inflates MSE and shrinks r_i . Instead use the variance estimate from the fit without observation i

$$S_{(i)}^2 = \frac{(n - p) \text{MSE} - e_i^2 / (1 - h_{ii})}{n - p - 1}, \quad (4)$$

which also needs no refit.

R - Student Residuals

$$t_i = \frac{e_i}{\sqrt{S_{(i)}^2 (1 - h_{ii})}}$$

- Replaces MSE in (3) by $S_{(i)}^2$ from (4)
- If the errors are normal, $t_i \sim t_{n-p-1}$
- More sensitive to an outlier than r_i : the outlier no longer inflates its own denominator
- Used to test for outliers (later lecture)

Types of Residuals

Residual	Formula	Variance / distribution	Note
Raw	e_i	$\sigma^2(1 - h_{ii})$	Unequal variances
Standardized	$e_i / \sqrt{\text{MSE}}$	≈ 1	Ignores h_{ii}
Studentized	$e_i / \sqrt{\text{MSE}(1 - h_{ii})}$	1	Accounts for leverage
PRESS	$e_i / (1 - h_{ii})$	$\sigma^2 / (1 - h_{ii})$	Leave-one-out error
R - student	$e_i / \sqrt{S_{(i)}^2(1 - h_{ii})}$	t_{n-p-1}	Outlier detection

Example: Delivery Time Data

Delivery Time Data

- y : delivery time (minutes) of a route driver servicing vending machines
- x_1 : number of cases of product stocked
- x_2 : distance walked by the driver (feet)
- $n = 25$ observations, $p = 3$

Fitted model

$$\hat{y} = 2.341 + 1.616 x_1 + 0.0144 x_2, \quad \text{MSE} = 10.624, \quad \hat{\sigma} = 3.259.$$

Delivery Time Data: Residuals

i	h_{ii}	e_i	d_i	r_i	$e_{(i)}$	t_i
1	0.102	-5.028	-1.543	-1.628	-5.598	-1.696
9	0.498	7.420	2.276	3.214	14.789	4.311
10	0.196	2.376	0.729	0.813	2.957	0.807
15	0.041	0.671	0.206	0.210	0.700	0.206
20	0.102	-5.788	-1.776	-1.874	-6.443	-1.997
22	0.392	-3.687	-1.131	-1.450	-6.059	-1.490

Six of the 25 observations. d_i standardized, r_i studentized, $e_{(i)}$ PRESS, t_i R-student. PRESS = 459.04.

Observation 9

Observation 9 (30 cases, 1460 ft) has the largest leverage, $h_{99} = 0.498$. Its residuals are $d_9 = 2.276$, $r_9 = 3.214$, $t_9 = 4.311$. Why does the residual keep growing from d_9 to r_9 to t_9 ?

Normality

Normality


Normal Probability Plot

Normal Probability Plot: Construction

- Order the residuals: $e_{[1]} \leq e_{[2]} \leq \dots \leq e_{[n]}$
- Plot $e_{[i]}$ against the normal quantile $\Phi^{-1}\left(\frac{i - 1/2}{n}\right)$
- If the errors are normal, the points fall close to a straight line
- Use studentized or R-student residuals so that all residuals have the same variance


Normal Probability Plot: Patterns

- Straight line: consistent with normality
- Heavy tails: ends bend away from the line (S-shape)
- Light tails: ends bend toward the center (reversed S)
- Skewness: one-sided curvature
- Isolated points at the ends: outliers



Normal Probability Plot: Delivery Time Data

- The central points follow the line
- Observation 9 ($r_9 = 3.21$) lies far above the line in the upper tail
- Is the departure from normality serious, or is it one point?



Normality

Shapiro–Wilk Test

Shapiro–Wilk Test

H_0 : the errors are normally distributed, against H_1 : they are not.

$$W = \frac{\left(\sum_{i=1}^n a_i e_{[i]}\right)^2}{\sum_{i=1}^n (e_i - \bar{e})^2}$$

- $e_{[i]}$ are the ordered residuals; the constants a_i come from the expected order statistics of a normal sample and are supplied by software
- W is closely related to the squared correlation between ordered residuals and normal scores
- $0 < W \leq 1$: values near 1 are consistent with normality, small W rejects H_0

Test or Plot?

With $n = 100,000$ the Shapiro–Wilk test rejects for negligible departures from normality; with $n = 10$ it rarely rejects. Which should we trust, the test or the plot?

Delivery Time Data: Shapiro–Wilk

Applied to the studentized residuals: $W = 0.923$, $p = 0.060$.

- Do not reject normality at the 5% level, but the evidence is borderline
- The plot located the problem: observation 9
- Refitting without observation 9 (diagnostic only): $W = 0.970$, $p = 0.68$

Constant Variance

Constant Variance

Residuals Against Fitted Values


Residuals Against Fitted Values

- Plot r_i (or e_i) against $\hat{y}_i$
- If the model is adequate, the points form a horizontal band around zero with no structure
- Any systematic pattern points to a violated assumption

Why $\hat{y}_i$ and not y_i ?


Pattern (a): Satisfactory

- Horizontal band centered at zero
- No trend, no change in spread
- No evidence against linearity or constant variance




Pattern (b): Funnel

- Spread of the residuals grows with $\hat{y}_i$
- The variance is not constant, typically $\text{Var}(y)$ increases with $E(y)$
- Remedies:
variance-stabilizing transformation of y ,
weighted least squares




Pattern (c): Double Bow

- Spread is largest in the middle and smallest at both ends
- Typical when y is a proportion in $(0, 1)$, where $\text{Var}(y) \propto p(1 - p)$
- Remedy: transformation of y



Delivery Time Data: Residuals Against $\hat{y}_i$


- Most points form a roughly horizontal band
- Observations 9 and 22 stand apart at large fitted values
- Is the spread at large $\hat{y}_i$ real, or driven by these two points?




Name the Violation

Which assumption does each plot call into question?

(a)



(c)



Constant Variance

Breusch–Pagan and White Tests

Why a Test?

A funnel plot is a judgment call. Formal tests of

$$H_0 : \text{Var}(\varepsilon_i) = \sigma^2 \text{ for all } i$$

H_1 : the variance depends on the regressors

Two common choices: Breusch–Pagan and White.

Breusch-Pagan Test

- ① Fit the model and obtain the residuals e_i
- ② Regress e_i^2 on the k regressors and record R_{aux}^2
- ③ Compute $LM = n R_{\text{aux}}^2$
- ④ Under H_0 , $LM \sim \chi_k^2$ approximately for large n ; reject for large LM

This is the studentized form, the default in `bptest` in R.

White Test

Same steps as Breusch–Pagan, but the auxiliary regression of e_i^2 uses the regressors, their squares and their cross-products. For $k = 2$ regressors

$$e_i^2 \text{ on } x_1, x_2, x_1^2, x_2^2, x_1x_2 \implies LM \sim \chi_5^2.$$

- No assumption on the form of the heteroscedasticity
- Many auxiliary regressors when k is large, which lowers power

Delivery Time Data: Tests for Nonconstant Variance

Test	LM	df	p -value
Breusch–Pagan	11.99	2	0.0025
White	14.96	5	0.0105

- Both reject constant variance at the 5% level
- The rejection is driven by observation 9: refitting without it (diagnostic only) gives Breusch–Pagan $LM = 3.22$, $p = 0.20$

If the Variance Is Not Constant

- Variance-stabilizing transformation of y
- Weighted least squares (WLS)

Before either, check whether a single unusual observation, as in the delivery time data, produced the rejection.


Linearity of the Mean Function

Linearity of the Mean Function

Residuals Against Fitted Values

Pattern (d): Curvature


- Residuals are systematically positive in one range of $\hat{y}_i$ and negative in another
- The mean function is wrong: the linear model is inadequate
- Remedies: add higher-order or interaction terms, transform the regressors or y




Name the Violation

Which assumption does each plot call into question?

(i)



(ii)



Linearity of the Mean Function



Residuals Against Regressors

Residuals Against Regressors

Plot r_i against each regressor x_j .

- Curvature: x_j enters the mean function in a nonlinear way, add a higher-order term or transform x_j
- Funnel: the variance depends on x_j
- A plot against a regressor not in the model: a pattern means that regressor should be added

Delivery Time Data: Residuals Against x_1 and x_2



Neither plot shows clear curvature. Observation 9 is isolated at the extreme of both regressors.


Uncorrelated Errors

Uncorrelated Errors


Residuals Against Time Order

Residuals Against Time Order

Plot e_t (or r_t) against the order in which the data were collected.



Positive autocorrelation: long runs of one sign.



Negative autocorrelation: signs alternate.

A trend over time suggests a missing time-dependent regressor.

Uncorrelated Errors

Durbin-Watson Test

First-Order Autoregressive Errors

$$\varepsilon_t = \phi \varepsilon_{t-1} + a_t, \quad |\phi| < 1, \quad a_t \stackrel{iid}{\sim} N(0, \sigma_a^2)$$

$$E[\varepsilon_t] = 0$$

$$\text{Var}(\varepsilon_t) = \frac{\sigma_a^2}{1 - \phi^2}$$

$$\text{Corr}(\varepsilon_t, \varepsilon_{t+j}) = \phi^j$$

$\phi > 0$: neighboring errors are similar (positive autocorrelation). $\phi < 0$: they tend to alternate in sign.

Consequences of Autocorrelated Errors

- The least-squares estimators are still unbiased, but no longer have minimum variance
- MSE typically underestimates σ^2
- The standard errors of the coefficients are underestimated
- Confidence intervals and t , F tests are no longer valid

Durbin–Watson Statistic

$$d = \frac{\sum_{t=2}^n (e_t - e_{t-1})^2}{\sum_{t=1}^n e_t^2} \quad (5)$$
$$\approx 2(1 - r_1), \quad r_1 = \frac{\sum_{t=2}^n e_t e_{t-1}}{\sum_{t=1}^n e_t^2}$$

- r_1 is the lag - 1 sample autocorrelation of the residuals
- $0 \leq d \leq 4$
- $d \approx 2$: no autocorrelation; $d < 2$: positive; $d > 2$: negative

Why $d \approx 2(1 - r_1)$?

Start from the numerator of (5), $\sum_{t=2}^n (e_t - e_{t-1})^2$. Where does $2(1 - r_1)$ come from?

Decision Rule


Test $H_0 : \phi = 0$ against $H_1 : \phi > 0$. The bounds d_L and d_U depend on n , the number of regressors k and α , and are tabulated.

- $d < d_L$: reject H_0 , positive autocorrelation
- $d > d_U$: do not reject H_0
- $d_L \leq d \leq d_U$: inconclusive

For $H_1 : \phi < 0$ apply the same rule to $4 - d$, with d from (5).

Soft Drink Sales Data

- y : annual sales in a sales region, x : annual advertising expenditure, $n = 20$ years
- Fitted model:
 $\hat{y} = 1608.5 + 20.091 x$
- $d = 1.08$
- $\alpha = 0.05$, $n = 20$, $k = 1$:
 $d_L = 1.20$, $d_U = 1.41$
- $d < d_L$: reject H_0 , evidence of positive autocorrelation



The plot alone is hard to judge by eye.

If the Errors Are Autocorrelated

- First look for a missing regressor, often a time-related one
- Otherwise estimate the model with an iterative procedure such as Cochrane–Orcutt (MPV, Chapter 14)

Summary

Diagnostics at a Glance

Tool	Checks	Typical remedy
Normal plot, Shapiro–Wilk	normal errors	transform y
Plot vs $\hat{y}_i$, Breusch–Pagan, White	constant variance	transform y , WLS
Plot vs $\hat{y}_i$ and vs x_j	mean function	add terms, transform
Plot vs time, Durbin–Watson	uncorrelated errors	add regressor
$r_i, t_i, e_{(i)}$	unusual observations	investigate

Workflow and Key Messages

- ① Fit the model and compute studentized and R-student residuals
- ② Plot: normal probability plot, against $\hat{y}_i$, against regressors, against time order
- ③ Use formal tests to back up what the plots suggest
- ④ Apply a remedy, refit, and check again
 - Residuals are not the errors: they have unequal variances and are correlated
 - Plots first, tests second; tests depend heavily on n
 - A single observation can drive a rejection

Questions?
