Residual Analysis

Table of Contents

1. Model Adequacy Checking
2. Raw Residuals
3. Different Type of Residuals
4. Example: Delivery Time Data
5. Normality

Table of contents continued in the next page...

Table of Contents (cont.)

6. Constant Variance
7. Linearity of the Mean Function
8. Uncorrelated Errors
9. Summary

Model Adequacy Checking

Assumptions Behind the Fitted Model

The linear model $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$ is fitted under the following assumptions:

- The relationship between y and the regressors is (approximately) linear
- $\mathrm{E}[\varepsilon_i] = 0$
- $\mathrm{Var}(\varepsilon_i) = \sigma^2$ for all i (constant variance)
- The errors are uncorrelated (independence)
- The errors are normally distributed (needed for the inference)

Least squares returns estimates whether or not these assumptions hold.

What Can We Observe?

The errors ε_i are never observed. What can we observe instead?

Raw Residuals

Raw Residuals

The ith residual is $e_i = y_i - \hat{y}_i$. In matrix form

$$\begin{aligned}\mathbf{e} = \mathbf{y} - \hat{\mathbf{y}} = \mathbf{y} - \mathbf{H}\mathbf{y} &= (\mathbf{I} - \mathbf{H})\mathbf{y} \\ &= (\mathbf{I} - \mathbf{H})(\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}) \\ &= (\mathbf{I} - \mathbf{H})\boldsymbol{\varepsilon},\end{aligned}$$

where the last step uses $(\mathbf{I} - \mathbf{H})\mathbf{X} = \mathbf{0}$.

Leverage

The hat matrix is $\mathbf{H} = \mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'$ with diagonal elements

$$h_{ii} = \mathbf{x}_i'(\mathbf{X}'\mathbf{X})^{-1}\mathbf{x}_i.$$

- Intuitively, h_{ii} measures how far $\mathbf{x}_i$ is from the centroid of the regressor values
- h_{ii} is large when $\mathbf{x}_i$ is far from the centroid of the regressor values
- Points with unusual $\mathbf{x}_i$-values have more "leverage" to pull the fitted surface toward their own y_i
- $0 \leq h_{ii} \leq 1$ and $\sum_{i=1}^{n} h_{ii} = p$, where p is the number of model parameters
- $\hat{y}_i = \sum_j h_{ij} y_j$, so h_{ii} is the weight of y_i in its own fitted value

Flagging High-Leverage Points

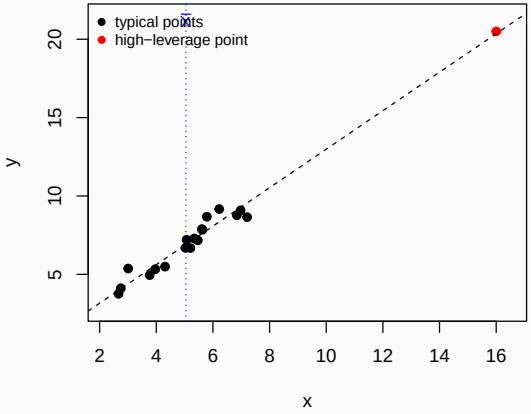
Rule of thumb: observation i has high leverage if

$$h_{ii} > \frac{2p}{n}.$$

- Delivery time data: $p = 3$, $n = 25$, so the cutoff is $2p/n = 2(3)/25 = 0.24$
- Observation 9: $h_{99} = 0.498 > 0.24$, flagged as high leverage
- High leverage is about an unusual $\mathbf{x}_i$, not an unusual y_i: a flagged point can still have a small residual if y_i happens to agree with the trend at that $\mathbf{x}_i$

Leverage: Geometric Intuition

- Leverage depends only on the location of $\mathbf{x}_i$, not on y_i
- A point far from $\bar{x}$ has high leverage, whatever its y-value
- Here the flagged point sits far from the rest of the data in x, yet close to the fitted line: high leverage, small residual



Mean and Covariance of the Residuals

$$\begin{aligned} \mathrm{E}[\mathbf{e}] &= \mathbf{0} \\ \mathrm{Var}(\mathbf{e}) &= \sigma^2(\mathbf{I} - \mathbf{H}) \end{aligned} \tag{1}$$

Element by element

$$\begin{aligned} \mathrm{Var}(e_i) &= \sigma^2(1 - h_{ii}) \\ \mathrm{Cov}(e_i, e_j) &= -\sigma^2 h_{ij}, \qquad i \neq j. \end{aligned}$$

Are the Residuals iid?

The errors are assumed iid $N(0, \sigma^2)$. Are the residuals $e_1, \ldots, e_n$ iid?

No. $\mathrm{Var}(e_i) = \sigma^2(1 - h_{ii})$ differs across i whenever the h_{ii} differ, $\mathrm{Cov}(e_i, e_j) = -\sigma^2 h_{ij} \neq 0$, and the residuals satisfy the p linear constraints $\mathbf{X}'\mathbf{e} = \mathbf{0}$ (in particular $\sum_i e_i = 0$ when the model has an intercept).

Are the Residuals iid?

The errors are assumed iid $N(0, \sigma^2)$. Are the residuals $e_1, \ldots, e_n$ iid?

No. $\mathrm{Var}(e_i) = \sigma^2(1 - h_{ii})$ differs across i whenever the h_{ii} differ, $\mathrm{Cov}(e_i, e_j) = -\sigma^2 h_{ij} \neq 0$, and the residuals satisfy the p linear constraints $\mathbf{X}'\mathbf{e} = \mathbf{0}$ (in particular $\sum_i e_i = 0$ when the model has an intercept).

Different Type of Residuals

Standardized Residuals

$$d_i = \frac{e_i}{\sqrt{\mathrm{MS}_{\mathrm{Res}}}}, \qquad \mathrm{MS}_{\mathrm{Res}} = \frac{\mathrm{SS}_{\mathrm{Res}}}{n-p} \tag{2}$$

- $\mathrm{E}[d_i] = 0$ and $\mathrm{Var}(d_i) \approx 1$, and d_i is free of the units of y
- Rule of thumb: $|d_i| > 3$ marks a potential outlier
- d_i ignores that $\mathrm{Var}(e_i) = \sigma^2(1 - h_{ii})$ differs across observations, see (1)

Studentized Residuals

$$r_i = \frac{e_i}{\sqrt{\mathrm{MS}_{\mathrm{Res}}\,(1 - h_{ii})}} = \frac{d_i}{\sqrt{1 - h_{ii}}} \tag{3}$$

- Divides e_i by its own estimated standard deviation from (1), so $\mathrm{Var}(r_i) = 1$ for every i, whatever the location of $\mathbf{x}_i$
- $|r_i| \geq |d_i|$, with d_i from (2)
- $r_i \approx d_i$ when h_{ii} is small, and they differ noticeably for high-leverage points

PRESS Residual: Leave-One-Out Prediction Error

The PRESS residual is the error when y_i is predicted from the model fitted to the other $n-1$ observations

$$e_{(i)} = y_i - \hat{y}_{(i)} = \frac{e_i}{1-h_{ii}}.$$

- The identity avoids refitting the model n times
- A large h_{ii} inflates the gap between $e_{(i)}$ and e_i

Variance of PRESS Residuals and the PRESS Statistic

$$\mathrm{Var}(e_{(i)}) = \frac{\sigma^2}{1 - h_{ii}}$$

$$\mathrm{PRESS} = \sum_{i=1}^{n} e_{(i)}^2$$

- PRESS measures predictive ability: smaller is better
- Standardizing $e_{(i)}$ by its standard deviation returns the studentized residual

$$\frac{e_{(i)}}{\sqrt{\mathrm{Var}(e_{(i)})}} = \frac{e_i}{\sqrt{\sigma^2(1 - h_{ii})}}$$

Removing Observation i From the Variance Estimate

MSE uses all n observations, including observation i: a large e_i inflates MSE and shrinks r_i. Instead use the variance estimate from the fit without observation i

$$S^2_{(i)} = \frac{(n-p)\,\mathrm{MS}_{\mathrm{Res}} - e_i^2/(1-h_{ii})}{n-p-1}, \tag{4}$$

which also needs no refit.

R-Student Residuals (cont.)

$$t_i = \frac{e_i}{\sqrt{S^2_{(i)}\,(1 - h_{ii})}}$$

- Replaces MSE in (3) by $S^2_{(i)}$ from (4)
- If the errors are normal, $t_i \sim t_{n-p-1}$
- More sensitive to an outlier than r_i: the outlier no longer inflates its own denominator
- Used to test for outliers (later lecture)

Types of Residuals

Residual	Formula	Variance / distribution	Note
Raw	e_i	$\sigma^2(1-h_{ii})$	Unequal variances
Standardized	$e_i/\sqrt{\mathrm{MS}_{\mathrm{Res}}}$	≈ 1	Ignores h_{ii}
Studentized	$e_i/\sqrt{\mathrm{MS}_{\mathrm{Res}}(1-h_{ii})}$	1	Accounts for leverage
PRESS	$e_i/(1-h_{ii})$	$\sigma^2/(1-h_{ii})$	Leave-one-out error
R-student	$e_i/\sqrt{S^2_{(i)}(1-h_{ii})}$	t_{n-p-1}	Outlier detection

Leverage and Residual Size

Points with large leverage tend to have small raw residuals. Why?

A large h_{ii} pulls the fitted surface toward y_i, since $\hat{y}_i = \sum_j h_{ij} y_j$ contains $h_{ii} y_i$. Formally $\mathrm{Var}(e_i) = \sigma^2(1 - h_{ii}) \to 0$ as $h_{ii} \to 1$, so e_i understates how far y_i is from where the rest of the data would put it.

Leverage and Residual Size

Points with large leverage tend to have small raw residuals. Why?

A large h_{ii} pulls the fitted surface toward y_i, since $\hat{y}_i = \sum_j h_{ij} y_j$ contains $h_{ii} y_i$. Formally $\mathrm{Var}(e_i) = \sigma^2(1 - h_{ii}) \to 0$ as $h_{ii} \to 1$, so e_i understates how far y_i is from where the rest of the data would put it.

Example: Delivery Time Data

Delivery Time Data

- y: delivery time (minutes) of a route driver servicing vending machines
- x_1: number of cases of product stocked
- x_2: distance walked by the driver (feet)
- $n = 25$ observations, $p = 3$

Fitted model

$$\hat{y} = 2.341 + 1.616\,x_1 + 0.0144\,x_2, \qquad \mathrm{MS}_{\mathrm{Res}} = 10.624, \quad \hat{\sigma} = 3.259.$$

Delivery Time Data: Residuals

i	h_{ii}	e_i	d_i	r_i	$e_{(i)}$	t_i
1	0.102	-5.028	-1.543	-1.628	-5.598	-1.696
9	0.498	7.420	2.276	3.214	14.789	4.311
10	0.196	2.376	0.729	0.813	2.957	0.807
15	0.041	0.671	0.206	0.210	0.700	0.206
20	0.102	-5.788	-1.776	-1.874	-6.443	-1.997
22	0.392	-3.687	-1.131	-1.450	-6.059	-1.490

Six of the 25 observations. d_i standardized, r_i studentized, $e_{(i)}$ PRESS, t_i R-student. $\mathrm{PRESS} = 459.04$.

Observation 9

Observation 9 (30 cases, 1460 ft) has the largest leverage, $h_{99} = 0.498$. Its residuals are $d_9 = 2.276$, $r_9 = 3.214$, $t_9 = 4.311$. Why does the residual keep growing from d_9 to r_9 to t_9?

$r_9 = d_9/\sqrt{1 - h_{99}} = d_9/0.708$: high leverage shrinks $\mathrm{Var}(e_9)$, so the raw residual understates how unusual the point is. t_9 replaces $\mathrm{MS}_{\mathrm{Res}} = 10.624$ by $S^2_{(9)} = 5.905$, the variance estimate without observation 9 itself: the outlier no longer inflates the denominator.

Observation 9

Observation 9 (30 cases, 1460 ft) has the largest leverage, $h_{99} = 0.498$. Its residuals are $d_9 = 2.276$, $r_9 = 3.214$, $t_9 = 4.311$. Why does the residual keep growing from d_9 to r_9 to t_9?

$r_9 = d_9/\sqrt{1 - h_{99}} = d_9/0.708$: high leverage shrinks $\mathrm{Var}(e_9)$, so the raw residual understates how unusual the point is. t_9 replaces $\mathrm{MS}_{\mathrm{Res}} = 10.624$ by $S^2_{(9)} = 5.905$, the variance estimate without observation 9 itself: the outlier no longer inflates the denominator.

Normality

Normality

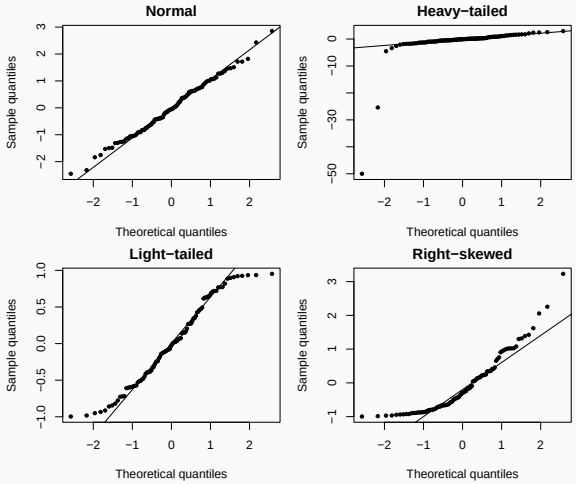
Normal Probability Plot

Normal Probability Plot: Construction

- Order the residuals: $e_{[1]} \leq e_{[2]} \leq ... \leq e_{[n]}$
- Plot $e_{[i]}$ against the normal quantile $\Phi^{-1}\left(\frac{i-1/2}{n}\right)$
- If the errors are normal, the points fall close to a straight line
- Use studentized or R-student residuals so that all residuals have the same variance

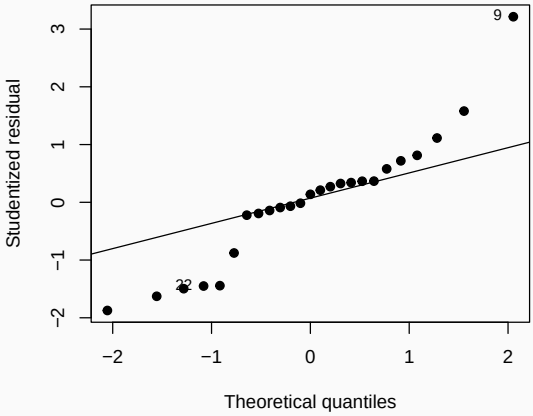
Normal Probability Plot: Patterns

- Straight line: consistent with normality
- Heavy tails: ends bend away from the line (S-shape)
- Light tails: ends bend toward the center (reversed S)
- Skewness: one-sided curvature
- Isolated points at the ends: outliers



Normal Probability Plot: Delivery Time Data

- The central points follow the line
- Observation 9 ($r_9 = 3.21$) lies far above the line in the upper tail
- Is the departure from normality serious, or is it one point?



Normality

Shapiro–Wilk Test

Shapiro–Wilk Test

H_0: the errors are normally distributed, against H_1: they are not.

$$W = \frac{\left(\sum_{i=1}^{n} a_i \, e_{[i]}\right)^2}{\sum_{i=1}^{n} (e_i - \bar{e})^2}$$

- $e_{[i]}$ are the ordered residuals; the constants a_i come from the expected order statistics of a normal sample and are supplied by software
- W is closely related to the squared correlation between ordered residuals and normal scores
- $0 < W \leq 1$: values near 1 are consistent with normality, small W rejects H_0

Test or Plot?

With $n = 100{,}000$ the Shapiro–Wilk test rejects for negligible departures from normality; with $n = 10$ it rarely rejects. Which should we trust, the test or the plot?

Delivery Time Data: Shapiro–Wilk

Applied to the studentized residuals: $W = 0.923$, $p = 0.060$.

- Do not reject normality at the 5% level, but the evidence is borderline
- The plot located the problem: observation 9
- Refitting without observation 9 (diagnostic only): $W = 0.970$, $p = 0.68$

Constant Variance

Constant Variance

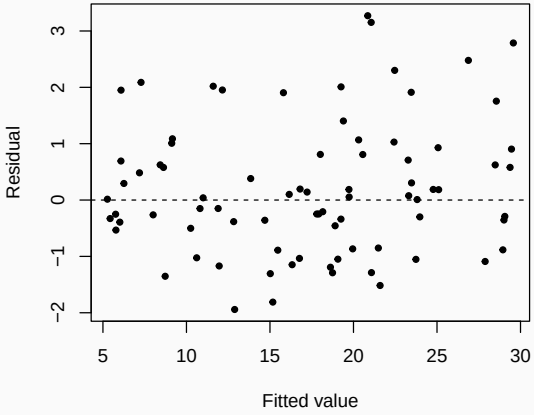
Residuals Against Fitted Values

Residuals Against Fitted Values

- Plot r_i (or e_i) against $\hat{y}_i$
- If the model is adequate, the points form a horizontal band around zero with no structure
- Any systematic pattern points to a violated assumption

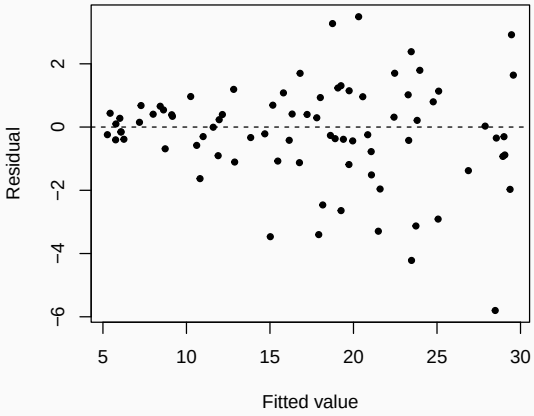
Pattern (a): Satisfactory

- Horizontal band centered at zero
- No trend, no change in spread
- No evidence against linearity or constant variance



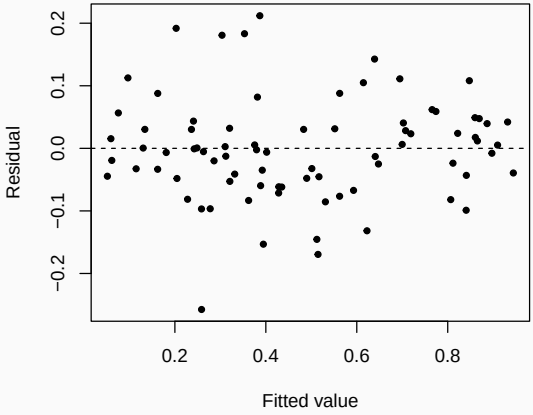
Pattern (b): Funnel

- Spread of the residuals grows with $\hat{y}_i$
- The variance is not constant, typically $\mathrm{Var}(y)$ increases with $\mathrm{E}[y]$
- Remedies: variance-stabilizing transformation of y, weighted least squares



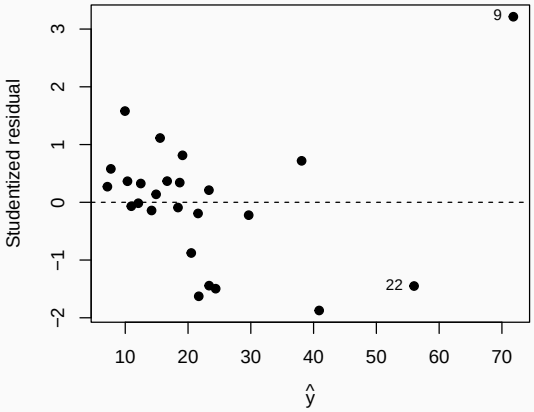
Pattern (c): Double Bow

- Spread is largest in the middle and smallest at both ends
- Typical when y is a proportion in $(0, 1)$, where $\mathrm{Var}(y) \propto p(1-p)$
- Remedy: transformation of y



Delivery Time Data: Residuals Against $\widehat{y}_i$

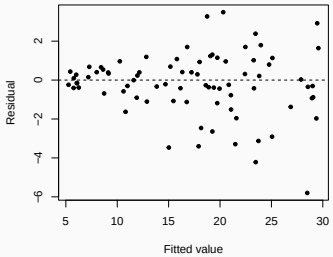
- Most points form a roughly horizontal band
- Observations 9 and 22 stand apart at large fitted values
- Is the spread at large $\widehat{y}_i$ real, or driven by these two points?



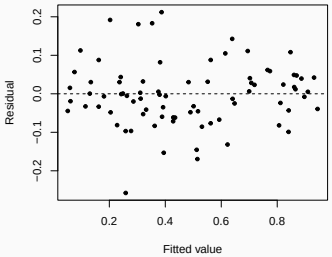
Name the Violation

Which assumption does each plot call into question?

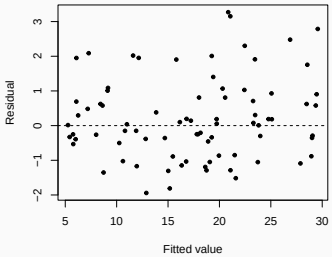
(i)



(ii)



(iii)



Constant Variance

Breusch–Pagan and White Tests

Why a Test?

A funnel plot is a judgment call. Formal tests of

$$H_0 : \mathrm{Var}(\varepsilon_i) = \sigma^2 \text{ for all } i$$
$$H_1 : \text{the variance depends on the regressors}$$

Two common choices: Breusch–Pagan and White.

Breusch–Pagan Test

1. Fit the model and obtain the residuals e_i
2. Regress e_i^2 on the k regressors and record R^2_{aux}
3. Compute $\text{LM} = n\, R^2_{\text{aux}}$
4. Under H_0, $\text{LM} \sim \chi^2_k$ approximately for large n; reject for large LM

This is the studentized form, the default in `bptest` in R.

White Test

Same steps as Breusch–Pagan, but the auxiliary regression of e_i^2 uses the regressors, their squares and their cross-products. For $k = 2$ regressors

$$e_i^2 \text{ on } x_1,\ x_2,\ x_1^2,\ x_2^2,\ x_1 x_2 \quad \Longrightarrow \quad \mathrm{LM} \sim \chi^2_5.$$

- No assumption on the form of the heteroscedasticity
- Many auxiliary regressors when k is large, which lowers power

Delivery Time Data: Tests for Nonconstant Variance

Test	LM	df	p-value
Breusch–Pagan	11.99	2	0.0025
White	14.96	5	0.0105

- Both reject constant variance at the 5% level
- The rejection is driven by observation 9: refitting without it (diagnostic only) gives Breusch–Pagan $\mathrm{LM} = 3.22$, $p = 0.20$

If the Variance Is Not Constant

- Variance-stabilizing transformation of y
- Weighted least squares (WLS)

Before either, check whether a single unusual observation, as in the delivery time data, produced the rejection.

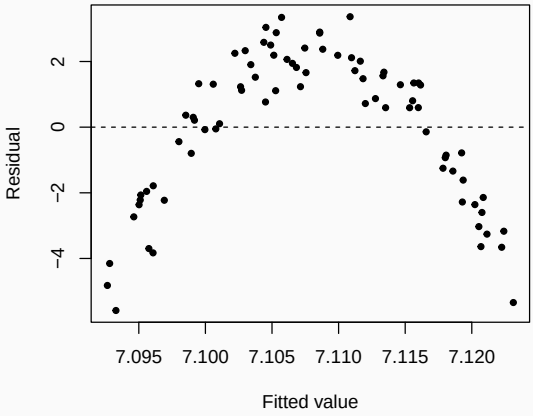
Linearity of the Mean Function

Linearity of the Mean Function

Residuals Against Fitted Values

Pattern (d): Curvature

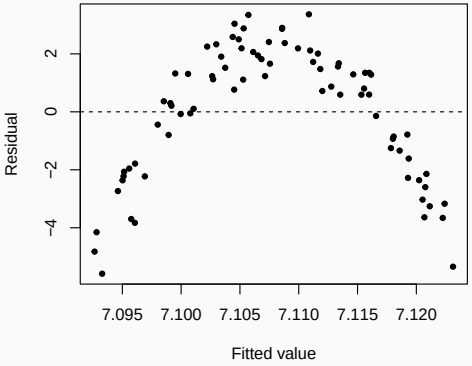
- Residuals are systematically positive in one range of $\hat{y}_i$ and negative in another
- The mean function is wrong: the linear model is inadequate
- Remedies: add higher-order or interaction terms, transform the regressors or y



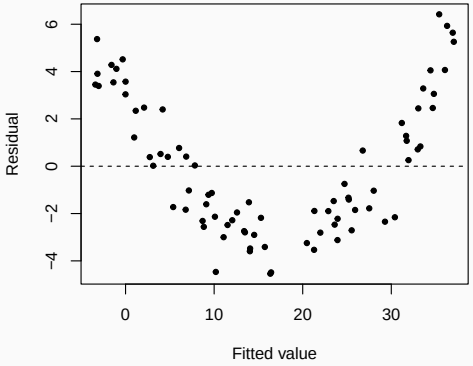
Name the Violation

Which assumption does each plot call into question?

(i)



(ii)



Linearity of the Mean Function

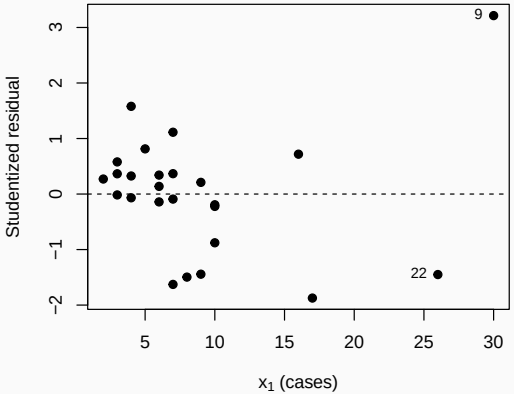
Residuals Against Regressors

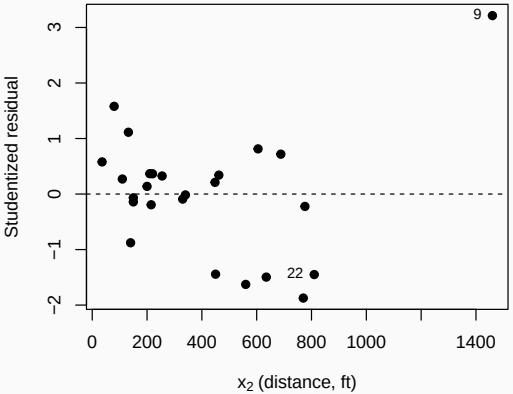
Residuals Against Regressors

Plot r_i against each regressor x_j.

- Curvature: x_j enters the mean function in a nonlinear way, add a higher-order term or transform x_j
- Funnel: the variance depends on x_j
- A plot against a regressor not in the model: a pattern means that regressor should be added

Delivery Time Data: Residuals Against x_1 and x_2





Neither plot shows clear curvature. Observation 9 is isolated at the extreme of both regressors.

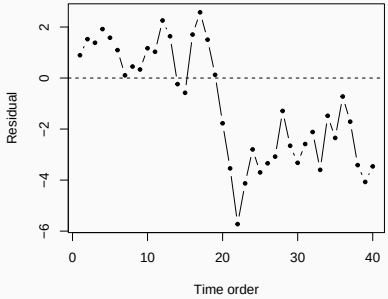
Uncorrelated Errors

Uncorrelated Errors

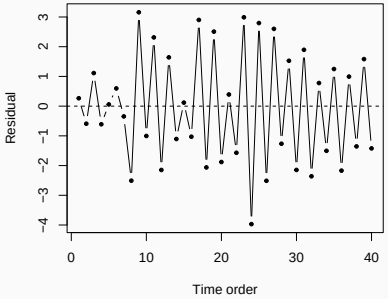
Residuals Against Time Order

Residuals Against Time Order

Plot e_t (or r_t) against the order in which the data were collected.



Positive autocorrelation: long runs of one sign.



Negative autocorrelation: signs alternate.

A trend over time suggests a missing time-dependent regressor.

Uncorrelated Errors

Durbin–Watson Test

First-Order Autoregressive Errors

$$\varepsilon_t = \phi\,\varepsilon_{t-1} + a_t, \qquad |\phi| < 1, \qquad a_t \overset{iid}{\sim} N(0, \sigma_a^2)$$

$$\begin{aligned} \mathrm{E}[\varepsilon_t] &= 0 \\ \mathrm{Var}(\varepsilon_t) &= \frac{\sigma_a^2}{1-\phi^2} \\ \mathrm{Corr}(\varepsilon_t, \varepsilon_{t+j}) &= \phi^j \end{aligned}$$

$\phi > 0$: neighboring errors are similar (positive autocorrelation). $\phi < 0$: they tend to alternate in sign.

Consequences of Autocorrelated Errors

- The least-squares estimators are still unbiased, but no longer have minimum variance
- MSE typically underestimates σ^2
- The standard errors of the coefficients are underestimated
- Confidence intervals and t, F tests are no longer valid

Durbin–Watson Statistic

$$d = \frac{\sum_{t=2}^{n} (e_t - e_{t-1})^2}{\sum_{t=1}^{n} e_t^2} \approx 2(1 - r_1), \qquad r_1 = \frac{\sum_{t=2}^{n} e_t e_{t-1}}{\sum_{t=1}^{n} e_t^2} \tag{5}$$

- r_1 is the lag-1 sample autocorrelation of the residuals
- $0 \leq d \leq 4$
- $d \approx 2$: no autocorrelation; $d < 2$: positive; $d > 2$: negative

Decision Rule

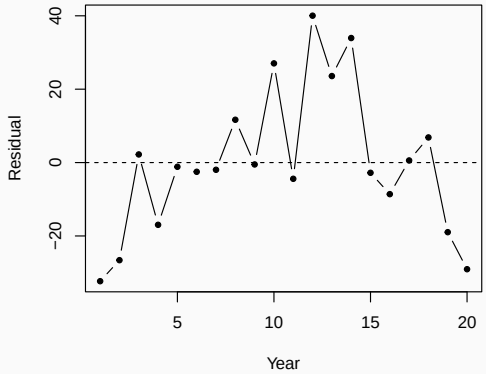
Test $H_0 : \phi = 0$ against $H_1 : \phi > 0$. The bounds d_L and d_U depend on n, the number of regressors k and α, and are tabulated.

- $d < d_L$: reject H_0, positive autocorrelation
- $d > d_U$: do not reject H_0
- $d_L \leq d \leq d_U$: inconclusive

For $H_1 : \phi < 0$ apply the same rule to $4 - d$, with d from (5).

Soft Drink Sales Data

- y: annual sales in a sales region, x: annual advertising expenditure, $n = 20$ years
- Fitted model: $\hat{y} = 1608.5 + 20.091\, x$
- $d = 1.08$
- $\alpha = 0.05$, $n = 20$, $k = 1$: $d_L = 1.20$, $d_U = 1.41$
- $d < d_L$: reject H_0, evidence of positive autocorrelation



The plot alone is hard to judge by eye.

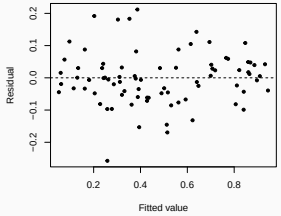
If the Errors Are Autocorrelated

- First look for a missing regressor, often a time-related one
- Otherwise estimate the model with an iterative procedure such as Cochrane–Orcutt (MPV, Chapter 14)

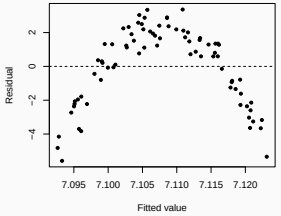
Summary

Review: Which Assumption Is Violated?

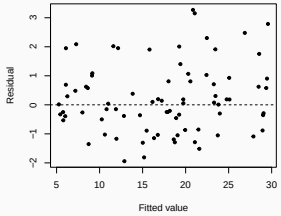
(i)



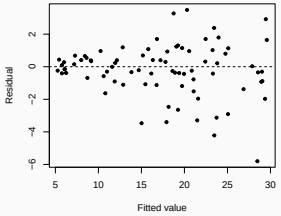
(ii)



(iii)



(iv)



Diagnostics at a Glance

Tool	Checks	Typical remedy
Normal plot, Shapiro–Wilk	Normal errors	Transform y
Plot vs $\hat{y}_i$, Breusch–Pagan, White	Constant variance	Transform y, WLS
Plot vs $\hat{y}_i$ and vs x_j	Mean function	Add terms, transform
Plot vs time, Durbin–Watson	Uncorrelated errors	Add regressor
r_i, t_i, $e_{(i)}$	Unusual observations	Investigate

Workflow and Key Messages

1. Fit the model and compute studentized and R-student residuals
2. Plot: normal probability plot, against $\hat{y}_i$, against regressors, against time order
3. Use formal tests to back up what the plots suggest
4. Apply a remedy, refit, and check again

- Residuals are not the errors: they have unequal variances and are correlated
- Plots first, tests second; tests depend heavily on n

Questions?